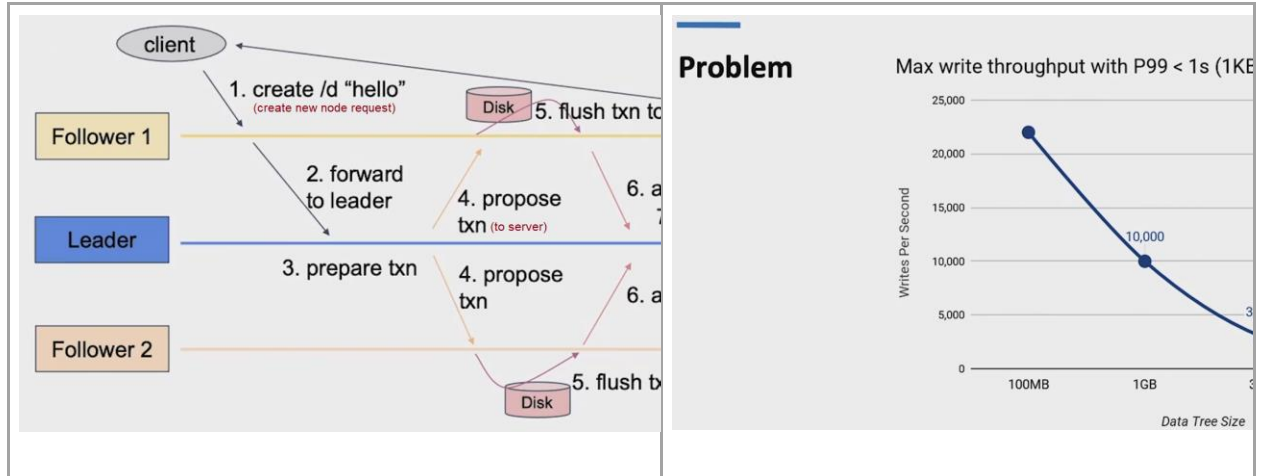


Zookeeper @ Facebook

- Facebook runs on ZK today and has a dedicated ZK team (8 engineers as of 01/2020) that contributes to hundreds of ZK use cases across FB engineering domains.
- ZK started @ FB circa 2014 for initial canonical use case of configuration & metadata management along with service discovery. Continued & expanded usage from then onwards with heavy current usage.
- ZK also being used for distributed semaphores & locks, membership management, and leader election @ FB
- Facebook running on 3.6 (FB internal ZK stack revision) in master for more than 2 years due to ZK enhancements (Peter considers that a notable achievement and vote of confidence in the stability of that branch)
- ZK engineering team continues to upstream new features, bug fixes, and enhancements around performance and reliability to its internal ZK stack.
- Recent updates on ZK @ FB (01/2020):
 - TLS enabled for all client connections (major effort due to broad distribution of clients)
 - ZK transitioned from bare metal to FB's container platform
 - Dramatic scale & efficiency improvements
 - Upstreamed almost all features (excluding C++ facebook ZK client) in version 3.6
 - 12 inconsistency & data corruption bugs fixed as of 01/2020
 - 120 patches upstreamed between 2019-2020
- Key ZK features slated for upstream in FY20:
 - Inspect API
 - Continuously monitors performance bottlenecks across ZK stack
 - Snapshot scheduler
 - Reduces disk I/O contention in ZK and expands maximum data traffic volume
 - Batch ACK and Commit
 - Extends overall max throughput thresholds in ZK
 - Failure request short circuit
 - Improves ZK quality & efficiency
- Additional ZK features slated for upstream:
 - Server-side semaphore API
 - Service isolation & demand control
 - Live traffic migration w/ ensemble split feature
 - Global session inconsistency detector
 - OOM (memory) protection

Facebook's Zookeeper Snapshot Scheduler Improvements:

Scheduler Write Flow (High I/O Latency Due to Disk Flush Concurrency):



Scheduler Write Flow Improvement (Simplified. Solves for High Latency via SnapPing Request/ACK):

